

# Databases, SQL, and Django

Scott Coughlin, Senior Computational Specialist

Some slides adapted from Professor Michael  
Coughlin and Matthew Graham

June 4<sup>th</sup> 2024

# Databases, SQL, and Django

## Topics Covered

1. Tabular Data
2. Servers, Databases, Schema, Relations, OH MY!
3. Doing Things the Structured Query Language (SQL) Way
4. Django
5. How to use docker-compose

# Tabular Data

# Tabular Data

```
==> IMDB-directors.csv <==  
director_id,first_name,last_  
name  
1,Todd,1  
2,Les,12 Poissons  
3,Lejaren,a'Hiller  
4,Nian,A
```

```
==> IMDB-movies.csv <==  
movie_id,name,year,rank  
0,#28,2002,0.0  
1,"#7 Train: An Immigrant  
Journey, The",2000,0.0  
2,$,1971,6.4  
3,"$1,000 Reward",1913,0.0
```

```
==> IMDB-movies_directors.csv <==  
director_id,movie_id  
1,378879  
2,281325  
3,30621  
3,304743
```

```
==> IMDB-movies_genres.csv <==  
movie_id,genre  
1,Documentary  
1,Short  
2,Comedy  
2,Crime
```

## Tabular Data

What if you had something that could not only manage all of those tables, i.e. adding new rows, but also could manage the way that those tables relate to one another in a hyper efficient way?

That is where a Database Management System comes in handy.

# Some Popular Database Management System

- Different DBMS:
  - Microsoft/Sybase
  - MySQL
  - Oracle
  - PostgreSQL
  - Redis, Hadoop

# Servers, Databases, Schema, Relations, OH MY!

# Jargon For Days and Days and Days

- The Server
- The Database
- A Database Management System
- Relations (tables)
- Attributes (columns)
- Domain (constraints on values for a column)
- Schema

## What is the server?

It is the physical (or virtual) machine where the database (and the application managing the database lives).

- “The database runs at/lives at  
xxxx.ciera.northwestern.edu”

# What is a database?

- A structured collection of data residing on a computer system (server) that can be easily accessed, managed and updated.
- Data is organized according to a database model.
- A Database Management System (DBMS) is a software package designed to store and manage databases.

# Some Popular Database Management System

- Different DBMS:
  - Microsoft/Sybase
  - MySQL
  - Oracle
  - PostgreSQL
  - Redis, Hadoop

# Why use a Database Management System?

- Provides concurrent access
- data scalability, expandability and flexibility
- Security through managing access to what databases, tables, and even types of queries a individual user can make.
- efficient memory management and indexing of the data
- Integrity constraints

# How do we model our data? With tables (or relations)

- Data is organized as **relations (tables)**, **attributes (columns)** and **domains (type)**
- A **relation** is a table with columns (attributes) and rows (tuples)
- The **domain** is the set of values that the attributes are allowed to take
- Within the relation, each row is unique, the column order does not matter, and each row contains a single value for each of its columns

# Relational Example

## Relational Model

Activity Code	Activity Name
23	Patching
24	Overlay
25	Crack Sealing

Key = 24

Activity Code	Date	Route No.
24	01/12/01	I-95
24	02/08/01	I-66

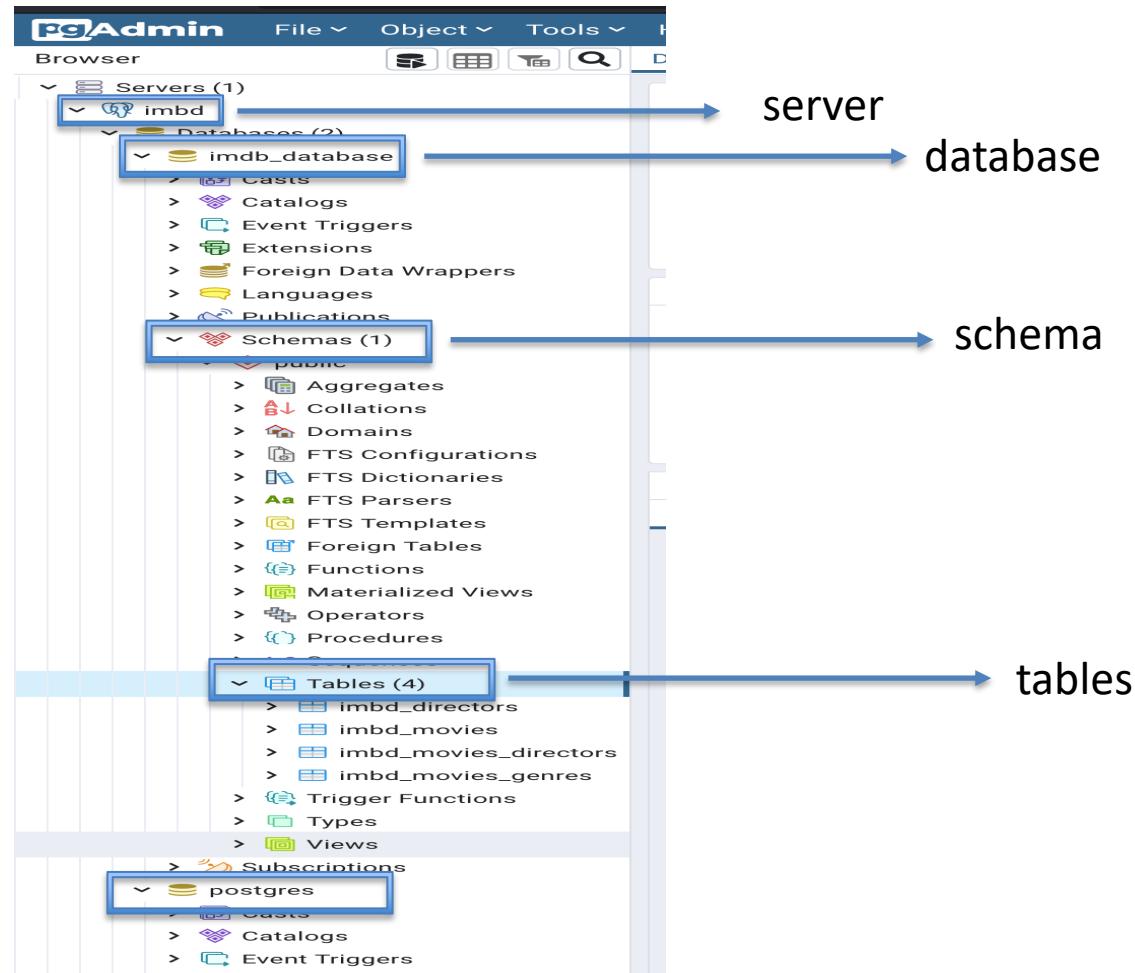
Date	Activity Code	Route No.
01/12/01	24	I-95
01/15/01	23	I-495
02/08/01	24	I-66

## What is the “Schema”

- The **schema** is a named collection of tables.
- A schema can contain many additional pieces of information relevant to a collection of tables including views, indexes, sequences, data types, operators, and functions.

# Say What Now?

# Say What Now?



# Doing things the Structured Query Language (SQL) Way

# Structured Query Language

- Different flavors:
  - DBMS – Flavor of SQL
  - Microsoft/Sybase - Transact-SQL
  - MySQL - MySQL
  - Oracle - PL/SQL
  - PostgreSQL - PL/pgSQL

# SELECT

```
SELECT column1, column2 FROM table WHERE  
condition (LIMIT #ofrows) ORDER BY  
sort_expression [ASC | DESC];
```

- ```
SELECT name, constellation FROM star WHERE dec > 0  
ORDER BY vmag;
```
- ```
SELECT * FROM star WHERE ra BETWEEN 0 AND 90;
```
- ```
SELECT DISTINCT constellation FROM star;
```
- ```
SELECT name FROM star LIMIT 5 ORDER BY vmag;
```

# JOIN

- Inner join: combining related rows

1. 

```
SELECT * FROM imbd_movies as s INNER JOIN imbd_movies_genres as t ON s.movie_id = t.movie_id;
```
2. 

```
SELECT * FROM star s, stellarTypes t WHERE s.stellarType = t.id;
```

- Outer join: each row does not need a matching row

1. 

```
SELECT * from star s LEFT OUTER JOIN stellarTypes t ON s.stellarType = t.id;
```
2. 

```
SELECT * from star s RIGHT OUTER JOIN stellarTypes t ON s.stellarType = t.id;
```
3. 

```
SELECT * from star s FULL OUTER JOIN stellarTypes t ON s.stellarType = t.id;
```

# Aggregate Functions

## COUNT, AVG, MIN, MAX, SUM

```
SELECT COUNT(*) FROM star;
```

```
SELECT AVG(vmag) FROM star;
```

```
SELECT stellarType, MIN(vmag), MAX(vmag) FROM star GROUP BY stellarType;
```

```
SELECT stellarType, AVG(vmag), COUNT(id) FROM star GROUP BY stellarType HAVING vmag > 14;
```

# Create

- CREATE DATABASE *databaseName*;
- CREATE TABLE *tableName* (*name1 type1*, *name2 type2*, ...);

```
CREATE TABLE star (name varchar(20), ra float, dec float, vmag float);
```

- Data types:
  - Boolean, bit, tinyint, smallint, int, bigint;
  - real/float, double, decimal;
  - char, varchar, text, binary, blob, longblob;
  - date, time, datetime, timestamp

```
CREATE TABLE star (name varchar(20) not null, ra float default 0, ...);
```

# KEYS

A **primary key** is a unique identifier for a row and is automatically not null

```
CREATE TABLE star (name varchar(20), ra float, dec float,  
vmag float, CONSTRAINT PRIMARY KEY (name));
```

A **foreign key** is a referential constraint between two tables identifying a column in one table that refers to a column in another table.

```
CREATE TABLE star (name varchar(20), ...,  
stellarType varchar(8), CONSTRAINT stellarType_fk  
FOREIGN KEY (stellarType) REFERENCES  
stellarTypes(id));
```

## Show and Describe

**SHOW ...**

SHOW INDEXES IN star;

SHOW WARNINGS;

**DESCRIBE...**

DESCRIBE star;

# INSERT

**INSERT INTO *table* VALUES(val1, val2, ...);**

```
INSERT INTO star VALUES('Sirius', 101.287, -16.716,  
-1.47);
```

```
INSERT INTO star(name, vmag) VALUES('Canopus', -  
0.72);
```

```
INSERT INTO star SELECT ...;
```

## DELETE

**DELETE FROM *table* WHERE *condition*;**

**DROP TABLE *table*;**

```
DELETE FROM star WHERE name = 'Canopus';
```

```
DELETE FROM star WHERE name LIKE 'C_n%';
```

# UPDATE

**UPDATE *table* SET *column* = *val1* WHERE *condition*;**

```
UPDATE star SET vmag = vmag + 0.5;
```

```
UPDATE star SET vmag = -1.47 WHERE name LIKE 'Sirius';
```

```
UPDATE star INNER JOIN temp on star.id = temp.id SET star.vmag =  
temp.mag;
```

# ALTER

**ALTER TABLE *table* ...;**

```
ALTER TABLE star ADD COLUMN bmag double AFTER vmag;
```

```
ALTER TABLE star DROP COLUMN bmag;
```

# Django

# Django

It seems like setting up and managing these databases would require learning a whole other language...can I not do that?

YES! Using a Python program like Django (or flask)

# Django

Do either of these print outs make you think of anything (from a Python programming perspective).

imdb_database=> SELECT * FROM imdb_movies LIMIT 2;				
movie_id	name	year	rank	
0	#28	2002	0	
1	"#7 Train: An Immigrant Journey, The"	2000	0	

```
(base) scottycoughlin@EDMRCSLC02G23LAMD6R SQL % head -n 3 IMDB-movies.csv  
movie_id,name,year,rank  
0,#28,2002,0.0  
1,"#7 Train: An Immigrant Journey, The",2000,0.0
```

## Some Hints

- Think about what a table is a bit more abstractly...
- It relates to something you are learning about this week...

# Django

The best way to think about a table, is from an *Object-Oriented Programming Perspective*.

The table is a **class** and the columns are its **attributes**.

```
class ImdbMovies(models.Model):
    movie_id = models.BigIntegerField(primary_key=True)
    name = models.TextField(blank=True, null=True)
    year = models.BigIntegerField(blank=True, null=True)
    rank = models.FloatField(blank=True, null=True)

    class Meta:
        db_table = 'imdb_movies'
```

# Django

This makes creating the relationships between the tables clearer as you can assign the values for one class as the column for a different class.

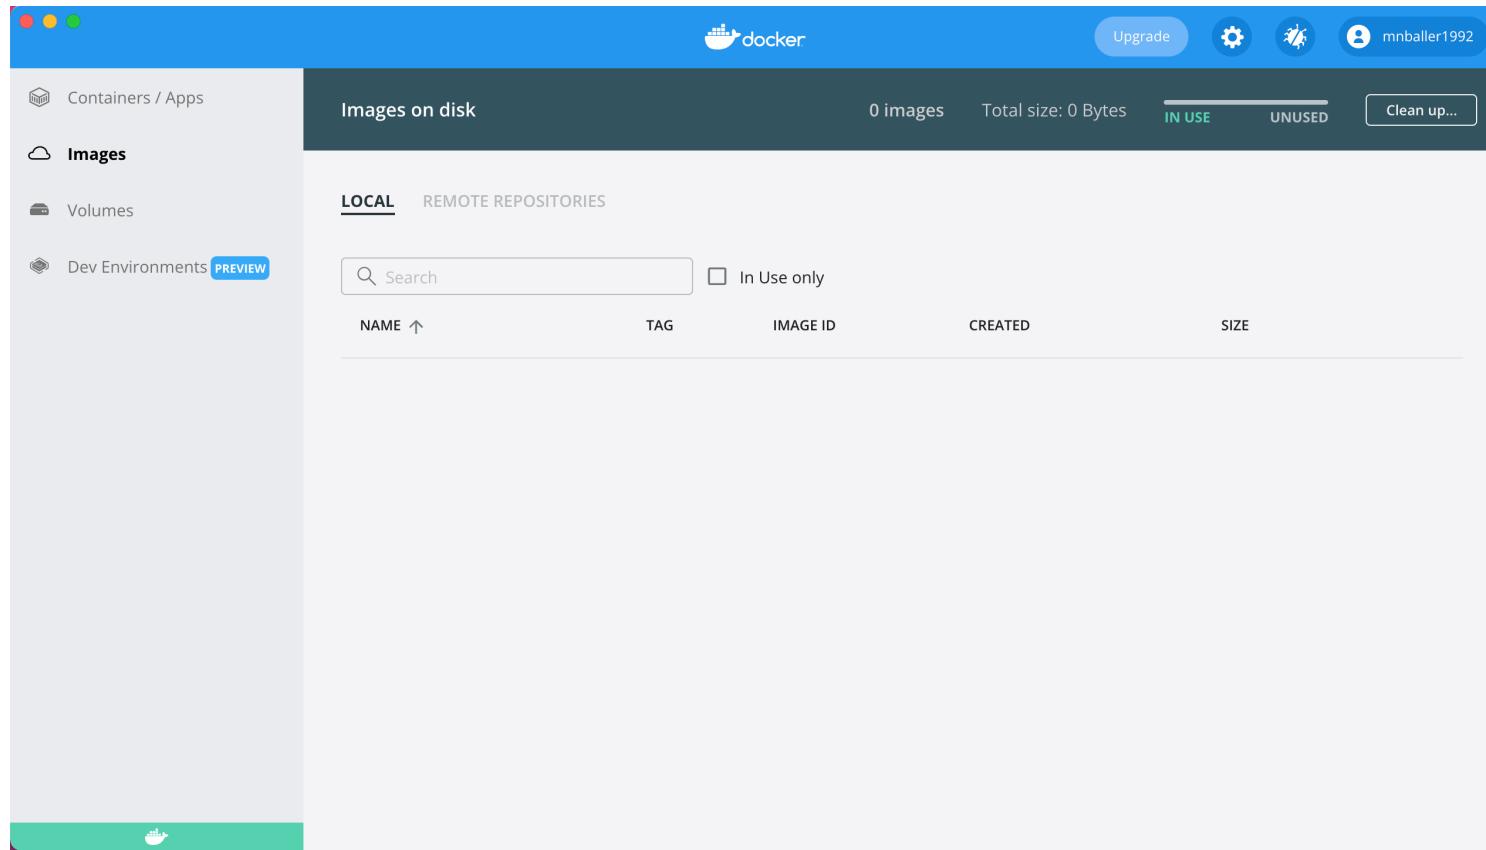
Below the column `movie_id` in the table `ImdbMoviesGenre` is the `ImdbMovies` class/table.

```
class ImdbMoviesGenre(models.Model):
    movie_id = models.ForeignKey(ImdbMovies, on_delete=models.CASCADE)
    genre = models.TextField(blank=True, null=True)

    class Meta:
        db_table = 'imdb_movies_genre'
```

# How to use docker-compose

# Launch DockerHub



# Open Terminal or PowerShell

```
(base) scottycoughlin@EDMRCSDL02G23LAMD6R SQL % cd ~/Documents/GitHub/Session-21/  
Day2/SQL/docker  
(base) scottycoughlin@EDMRCSDL02G23LAMD6R docker % cat README.md  
# Launch Docker Desktop  
In your applications folder, you can double click Docker Desktop and wait for it  
to launch  
  
# Launch Docker Containers  
From the current folder, i.e. `Session-21/Day2/SQL/docker`, run  
```  
docker-compose up --detach  
```
```

# docker-compose

```
# Mac or Linux  
$ cd Session-21/Day2/SQL/docker  
$ docker-compose up --detach
```

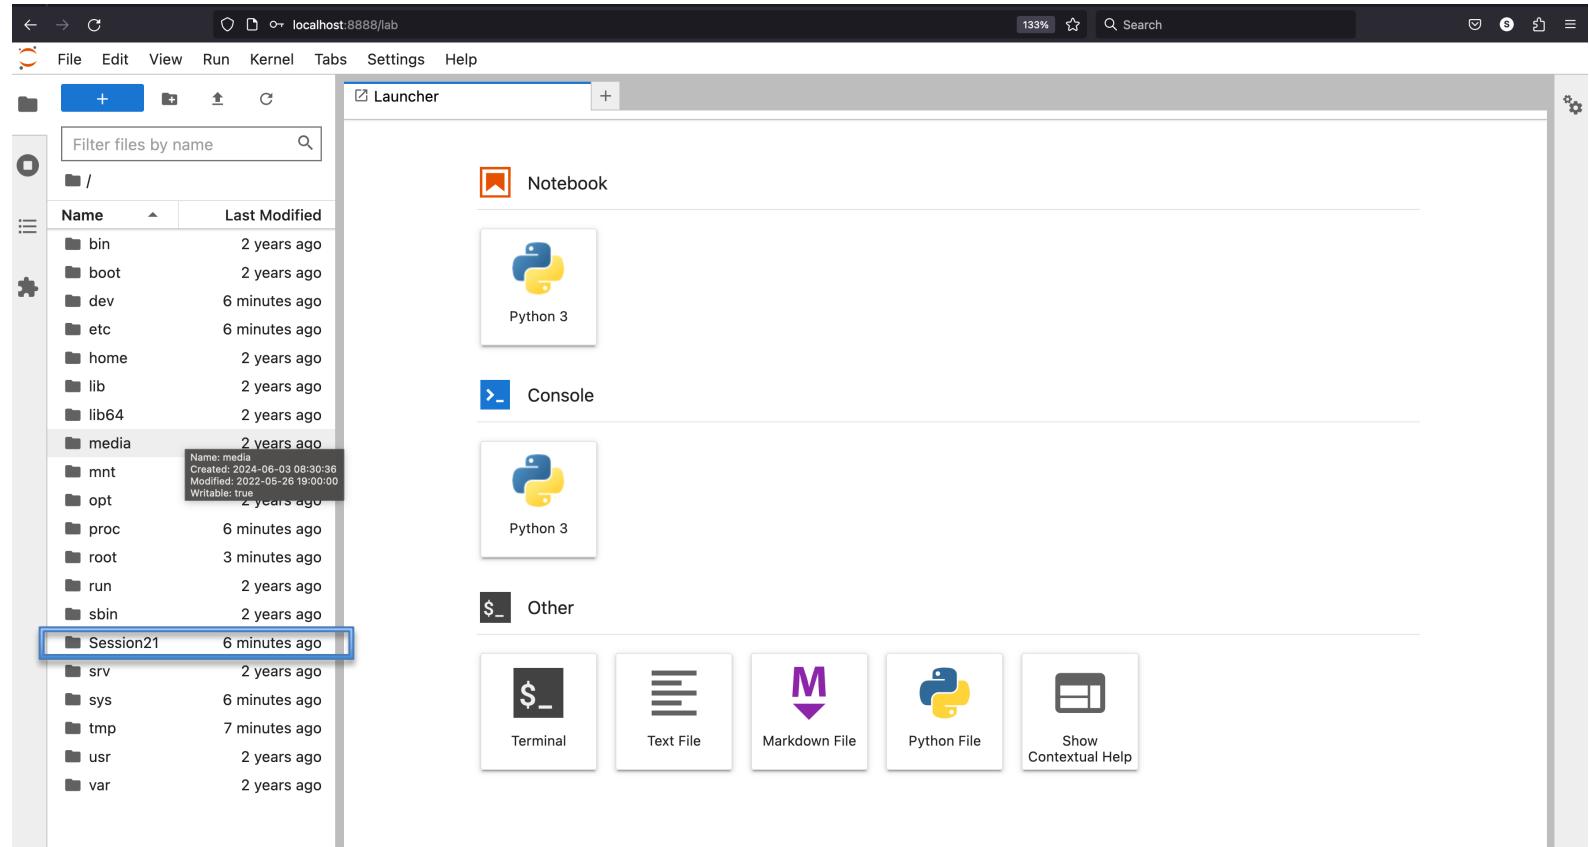
```
# Power Shell  
$ cd Session-21\Day2\SQL\docker  
$ docker-compose up --detach
```

- Three things are happening once you run this command!
  - A container with PostgreSQL installed is being downloaded, a database is being created and a table is being made for each “sheet” of IMBD data.
  - A useful web application, pgAdmin, is also being downloaded so that you can interact with Postgres via your browser.
  - A container with JupyterLab and Django is being made.

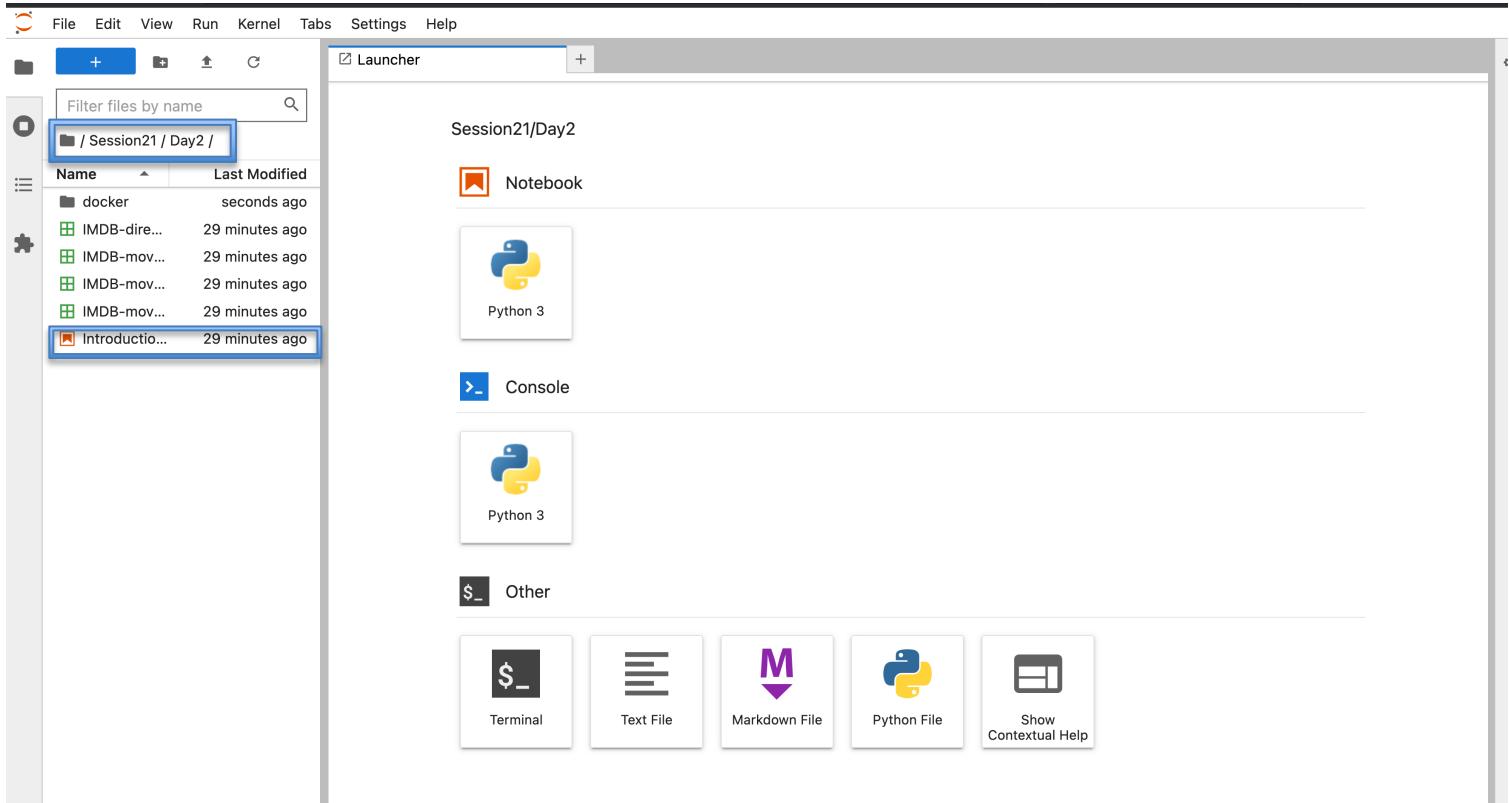
# JupyterLab

- Wait about 3 minutes and then...
- In your browser go to localhost:8888
- Password: Session21

# JupyterLab



# JupyterLab



# JupyterLab

The screenshot shows the JupyterLab interface running in a web browser. The top navigation bar includes File, Edit, View, Run, Kernel, Tabs, Settings, and Help. The tabs at the top show 'Launcher' (active), 'IntroductionToDjango.ipynb', and 'IntroductionToSQL.ipynb'. The left sidebar has a file browser with a tree view of files in '/Session21/Day2/'. The main content area contains a code editor with Python code related to SQLAlchemy and PostgreSQL, followed by a section titled 'Introduction to SQL (Structured Query Language)'.

```
[ ]: !cat /etc/os-release
[ ]: import numpy as np
[ ]: import pandas as pd
[ ]: import matplotlib.pyplot as plt

[ ]: from sqlalchemy import create_engine
[ ]: %matplotlib inline

[ ]: %load_ext sql
[ ]: %config SqlMagic.displaylimit=50
[ ]: %config SqMagic.autopandas=True

[ ]: %sql postgres://imdb:imdb_admin@postgres:5432/imdb_database

[ ]: connection = create_engine('postgresql://imdb:imdb_admin@postgres:5432/imdb_database')
```

## Introduction to SQL (Structured Query Language)

Version 0.1

By Scott Coughlin (Northwestern IT Research Computing and Data Services)  
June 4th 2024

Session 21 is primarily concerned with handling our data with efficiency.  
Ideally, for any and every task we want to desire solutions that operate *faster*.

This can be accomplished many different ways:

- build algorithms that execute faster
- spread calculations over many different computers simultaneously
- find a compact storage solution for the data so it can be accessed more quickly

In our introduction to SQL we will start with simple queries of existing tables, and discuss creating your own tables using `pandas` as a challenge problem.

### Problem 1) IMDb Data

# pgAdmin Portal

- In your browser go to [localhost:15432](http://localhost:15432)
- E-mail: [admin@pgadmin.com](mailto:admin@pgadmin.com)
- Password: password

# pgAdmin Portal

The screenshot shows the pgAdmin 4 web interface running at `localhost:15432/browser/`. The top navigation bar includes File, Object, Tools, Help, and a user account for `admin@pgadmin.com (internal)`. The main menu bar has tabs for Browser, Dashboard, Properties, SQL, Statistics, Dependencies, and Dependents. A sidebar on the left shows a tree view under Servers.

The central area features a "Welcome" section with the pgAdmin logo and the text "Management Tools for PostgreSQL". It highlights "Feature rich | Maximises PostgreSQL | Open Source". A note states: "You are currently running version 6.11 of pgAdmin 4, however the current version is 8.7. Please click [here](#) for more information." Below this, a "Quick Links" section contains a button labeled "Add New Server" with a server icon, and a "Configure pgAdmin" link with a gear icon.

The "Getting Started" section includes links to "PostgreSQL Documentation" (with a PostgreSQL icon), "pgAdmin Website" (with a globe icon), "Planet PostgreSQL" (with a document icon), and "Community Support" (with a people icon).

# pgAdmin Portal

Register - Server

General Connection SSL SSH Tunnel Advanced

Name: IMDB

Server group: Servers

Background: X

Foreground: X

Connect now?: On

Shared?: Off

Comments:

Register - Server

General Connection SSL SSH Tunnel Advanced

Host name/address: postgres

Port: 5432

Maintenance database: postgres

Username: imdb

Kerberos authentication?: Off

Password: ..... imdb\_admin

Save password?: Off

Role:

Service:

*?* *?* Close Reset Save

# pgAdmin Portal

The screenshot shows the pgAdmin 4 interface for the IMDB database. The left sidebar displays the database schema structure, including servers, databases, tables, and other objects. The main area features several performance monitoring dashboards:

- Database sessions:** Shows session counts (Total, Active, Idle) over time.
- Tuples in / Tuples out:** Shows the number of tuples inserted, updated, or deleted.
- Transactions per second:** A line chart showing transaction rates.
- Block I/O:** A line chart showing disk read and hit rates.
- Database activity:** A table showing session details like PID, User, Application, Client, Backend start, Transaction start, State, Wait event, and Blocking PIDs. One session is listed:

	PID	User	Application	Client	Backend start	Transaction start	State	Wait event	Blocking PIDs
X	■	▶ 133	imdb	pgAdmin 4 - DB:imdb_data...	172.18.0.3	2024-06-04 01:04:35 UTC	active		